

Problem Solving Question – Imaging Science

Assignment Report

Name: Bonthu Jayaram

1. Introduction

This assignment focuses on recovering the true texture (reflectance) of a surface from a single photograph that is unevenly illuminated. The idea is to separate the *actual surface pattern* from the *illumination effects* such as shadows, bright spots, or gradual light falloff.

The commonly used model for image formation is:

$$I(x, y) = R(x, y) L(x, y)$$

where

- $I(x,y)$ is the observed pixel intensity,
- $R(x,y)$ is the true reflectance (texture),
- $L(x,y)$ is the illumination (smooth lighting variations).

The challenge is that both R and L are unknown and only their product is visible.

2. Why Histogram Equalization Cannot Recover Reflectance

Histogram equalization is a global point-wise operation:

$$I_{\text{new}}(x, y) = f(I_{\text{old}}(x, y))$$

It treats all pixels only based on their intensity histogram, without understanding *why* a pixel is bright or dark.

It fails to recover $R(x, y)$ because:

1. No separation between lighting and texture

It cannot distinguish whether a pixel is dark due to actual texture or because it lies in a shadow.

2. No spatial or frequency awareness

Reflectance consists of **high-frequency details** (edges, patterns).

Illumination is **low-frequency** (smooth shading).

Histogram equalization ignores both properties.

3. Can distort texture

It often amplifies shadows or compresses high-frequency content, which corrupts the true surface texture.

Therefore, histogram equalization cannot extract reflectance. A frequency-based approach is needed.

3. Log-Domain Model for Reflectance Recovery

The multiplicative model can be converted to an additive model using logarithms:

$$\log I(x, y) = \log R(x, y) + \log L(x, y)$$

Let

- $i(x, y) = \log I(x, y)$
- $r(x, y) = \log R(x, y)$
- $l(x, y) = \log L(x, y)$

Then:

$$i(x, y) = r(x, y) + l(x, y)$$

Now, illumination becomes a **smooth, low-frequency component**, and texture becomes a **high-frequency component**.

This allows separation using simple smoothing and subtraction.

4. Proposed Algorithm (Retinex-style Method)

Step-by-step approach

1. Convert image to grayscale (for Question 2).
2. Normalize pixel values to $[0, 1]$.
3. Convert to log domain:

$$i(x, y) = \log (I(x, y) + \epsilon)$$

4. Estimate illumination $\hat{l}(x, y)$ by applying a **large manual box filter** (no built-ins).
5. Estimate reflectance:

$$\hat{r}(x, y) = i(x, y) - \hat{l}(x, y)$$

6. Convert back using exponent:

$$\hat{R}(x, y) = \exp(\hat{r}(x, y))$$

7. Normalize for visualization.

This produces a clear reflectance image with reduced shadows and preserved texture.

5. Grayscale Reflectance (Question 2)

Input Image (Your Photo)

(Insert your original grayscale or color image here)

Output Reflectance Image

(Insert your generated grayscale reflectance output here)

File: **reflectance_grayscale_output.jpg**

Observation

- The shadows cast by trees on the wall are significantly reduced.
- The brick pattern appears more uniform and clearer.
- Large-scale brightness variations are removed.
- Fine texture details of the bricks are preserved well.

This confirms successful separation of reflectance from illumination.

6. Color Reflectance (Question 3)

For a color image with spectral illumination variations, the model extends to RGB channels:

$$I_c(x, y) = R_c(x, y) L_c(x, y), c \in \{R, G, B\}$$

To maintain natural colors, a **single illumination estimate** is computed from the average log-intensity and applied uniformly across channels.

Output Reflectance Image (Color)

(Insert your generated color reflectance image here)

File: **reflectance_color_output.jpg**

Observation

- Joint illumination removal across all channels produced evenly balanced colors.
- True color ratios between bricks are preserved.

- Shadow regions no longer dominate the appearance.
 - The final output looks visually consistent and faithful to the material's real texture.
-

7. Conclusion

This assignment demonstrates how a single unevenly illuminated image can be decomposed into illumination and reflectance components. By using the logarithmic model and manual frequency separation:

- Illumination is treated as a slowly varying field.
- Reflectance retains fine texture information.
- Both grayscale and color outputs show effective shadow removal.

The results confirm that the Retinex-style approach is suitable for reflectance recovery even when only one image is available.

8. Appendix – Python Code

A. Grayscale Reflectance Code (Q2)

```
import numpy as np
from PIL import Image
import math

img = Image.open("C:/Users/Jayaram/Pictures/GREYSCALE.jpeg").convert("L")
I = np.asarray(img).astype(np.float32)

I_norm = I / 255.0
epsilon = 1e-6
log_I = np.log(I_norm + epsilon)

def manual_box_filter(image, kernel_size):
    h, w = image.shape
    k = kernel_size
```

```

pad = k // 2

padded = np.zeros((h + 2*pad, w + 2*pad), dtype=np.float32)
padded[pad:pad+h, pad:pad+w] = image

output = np.zeros_like(image)

area = k * k

for y in range(h):
    for x in range(w):
        window = padded[y:y+k, x:x+k]
        output[y, x] = np.sum(window) / area

return output

```

```

kernel_size = 41

log_L_hat = manual_box_filter(log_I, kernel_size)

log_R_hat = log_I - log_L_hat

```

```

R_hat = np.exp(log_R_hat)

R_hat = (R_hat - np.min(R_hat)) / (np.max(R_hat) + 1e-6)

R_hat_8bit = (R_hat * 255).astype(np.uint8)

```

```
Image.fromarray(R_hat_8bit).save("reflectance_grayscale_output.jpg")
```

B. Color Reflectance Code (Q3)

```

import numpy as np

from PIL import Image

img_color = Image.open("C:/Users/Jayaram/Pictures/GREYSCALE.jpeg").convert("RGB")

I_color = np.asarray(img_color).astype(np.float32)

R = I_color[:, :, 0] / 255.0

```

```
G = I_color[:, :, 1] / 255.0
```

```
B = I_color[:, :, 2] / 255.0
```

```
epsilon = 1e-6
```

```
log_R = np.log(R + epsilon)
```

```
log_G = np.log(G + epsilon)
```

```
log_B = np.log(B + epsilon)
```

```
log_intensity = (log_R + log_G + log_B) / 3.0
```

```
def manual_box_filter(image, kernel_size):
```

```
    h, w = image.shape
```

```
    k = kernel_size
```

```
    pad = k // 2
```

```
    padded = np.zeros((h + 2*pad, w + 2*pad), dtype=np.float32)
```

```
    padded[pad:pad+h, pad:pad+w] = image
```

```
    output = np.zeros_like(image)
```

```
    area = float(k*k)
```

```
    for y in range(h):
```

```
        for x in range(w):
```

```
            window = padded[y:y+k, x:x+k]
```

```
            output[y, x] = np.sum(window) / area
```

```
    return output
```

```
kernel_size = 41
```

```
log_L_hat = manual_box_filter(log_intensity, kernel_size)
```

```
log_Rr = log_R - log_L_hat
```

```
log_Rg = log_G - log_L_hat
```

```
log_Rb = log_B - log_L_hat
```

```
def normalize_channel(c):
```

```
    c = np.exp(c)
```

```
    c = c - np.min(c)
```

```
    c = c / (np.max(c) + 1e-6)
```

```
    return c
```

```
Rr = normalize_channel(log_Rr)
```

```
Rg = normalize_channel(log_Rg)
```

```
Rb = normalize_channel(log_Rb)
```

```
Rr8 = (Rr * 255).astype(np.uint8)
```

```
Rg8 = (Rg * 255).astype(np.uint8)
```

```
Rb8 = (Rb * 255).astype(np.uint8)
```

```
result_color = np.stack([Rr8, Rg8, Rb8], axis=2)
```

```
Image.fromarray(result_color).save("reflectance_color_output.jpg")
```

Input image



Grayscale image



Colour image



Question 2 – Computer Vision

1. Camera & Distortion Model + Robust Cost Function

1.1 Pinhole camera model

Let a 3D point on the grid plane be:

$$X = (X_w, Y_w, Z_w = 0, 1)^T$$

Camera projection without distortion:

1. Transform from world to camera coordinates:

$$\begin{bmatrix} X_c \\ Y_c \\ Z_c \end{bmatrix} = \begin{bmatrix} X_w \\ Y_w \\ 0 \end{bmatrix} + R \begin{bmatrix} Y_w \\ 0 \\ Z_w \end{bmatrix} + t$$

2. Normalize:

$$x = \frac{X_c}{Z_c}, y = \frac{Y_c}{Z_c}$$

3. Apply intrinsics:

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = K[y], K = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 1 & 1 & 0 \end{bmatrix}$$

where:

- f_x, f_y = focal lengths in pixels,
- c_x, c_y = principal point.

1.2 Radial distortion model (what we use)

We use a simple **two-parameter radial distortion** model:

$$\begin{aligned} r^2 &= x^2 + y^2 \\ x_d &= x(1 + k_1r^2 + k_2r^4) \\ y_d &= y(1 + k_1r^2 + k_2r^4) \end{aligned}$$

Then we apply intrinsics to **distorted** normalized coordinates:

$$\begin{aligned} u_d &= f_x x_d + c_x \\ v_d &= f_y y_d + c_y \end{aligned}$$

Here:

- k_1, k_2 are the radial distortion parameters we want to estimate.

So the full parameter vector is:

$$\theta = [f_x, f_y, c_x, c_y, k_1, k_2, R, t]$$

(we can represent R by 3 Rodrigues parameters; so 3 (rotation) + 3 (translation)).

1.3 Reprojection residuals

For each detected grid corner i :

- Known 2D image position: $p_i^{obs} = (u_i, v_i)$
- Known 3D plane point: $P_i = (X_{w,i}, Y_{w,i}, 0)$
- Predicted distorted image point from model: $p_i^{pred}(\theta) = (u_{d,i}(\theta), v_{d,i}(\theta))$

The **reprojection residual** is:

$$e_i(\theta) = \begin{bmatrix} u_i - u_{d,i}(\theta) \\ v_i - v_{d,i}(\theta) \end{bmatrix}$$

1.4 Robust cost function

Instead of minimizing plain sum of squared errors (which is sensitive to outliers), we use a **robust loss** (e.g., Huber):

$$E(\theta) = \sum_{i \in \mathcal{I}} \rho(\| e_i(\theta) \|^2)$$

Huber loss (concept):

$$\rho(s) = \begin{cases} s, & s \leq \delta^2 \\ 2\delta\sqrt{s} - \delta^2, & s > \delta^2 \end{cases}$$

- \mathcal{I} is the set of **inlier** points (from RANSAC).
- This cost function penalizes large errors sub-quadratically, making the optimization robust to remaining outliers.

This is the cost we conceptually optimize in our pipeline.

2. Robust Optimization Pipeline (Concept)

Pipeline overview:

1. **Detect grid corners** in the distorted image.
2. **Build 3D grid coordinates** on a plane ($Z = 0$, known spacing).
3. **RANSAC** to remove corner outliers.
4. **Non-linear optimization** (or library equivalent) to refine:
 - intrinsics K ,
 - radial distortion k_1, k_2 ,
 - extrinsics R, t .
5. **Undistort** the image and corner locations.
6. **Evaluate reprojection error**.

In code we will use OpenCV's `calibrateCamera`, which internally solves a very similar non-linear least squares problem (Levenberg–Marquardt). Our “robustness” comes from:

- Outlier rejection using **RANSAC** before calibration.
 - Restricting to a simple distortion model k_1, k_2 .
-

3. Using RANSAC to Remove Outliers

We have matches:

- 3D plane points $P_i = (X_{w,i}, Y_{w,i}, 0)$ (we know grid ordering),
- observed 2D points $p_i^{obs} = (u_i, v_i)$.

Idea:

- Ignore distortion for a moment and assume a **projective mapping** between plane and image: a **homography** H :

$$\begin{matrix} u_i \\ \lambda[v_i] \\ 1 \end{matrix} = H \begin{matrix} X_{w,i} \\ Y_{w,i} \\ 1 \end{matrix}$$

- Use **RANSAC homography estimation** to:
 - Randomly select minimal subsets of point matches.
 - Estimate H from each subset.
 - Count how many points are within a small reprojection threshold.
 - The model with the most inliers defines inlier set \mathcal{I} .

We then only use these **RANSAC inliers** for calibration.

In practice we'll use `cv2.findHomography(..., cv2.RANSAC, reprojThreshold)` which returns an inlier mask.

4. Undistort the Image & Compute Undistorted Grid

After we have estimated:

- Camera matrix K ,
- Distortion coefficients $d = [k_1, k_2, 0, 0, 0]$ (only radial used),

we can:

1. **Undistort full image**:
 - Compute an “optimized” new camera matrix K' using `cv2.getOptimalNewCameraMatrix`.

- Use cv2.undistort to map from distorted to undistorted image.

2. Undistort grid corners:

- Use cv2.undistortPoints to obtain undistorted positions of the detected corners.
- These form the **undistorted grid** on the ideal plane.

This undistorted grid is what the camera would see if the lens had no radial distortion.

5. Reproject Undistorted Grid & Compute Reprojection Error

To quantify how good our model is:

1. Take the **3D grid points** P_i (on the plane).
2. Use the estimated parameters (K, d, R, t) .
3. Project them back to the **distorted image** using the same forward model:
 - Use cv2.projectPoints with distortion coefficients.
4. Compare predicted positions p_i^{pred} to observed p_i^{obs} .

Compute the **reprojection error** per point:

$$e_i = \| p_i^{obs} - p_i^{pred} \|_2$$

We usually summarize with:

- Mean error:

$$\text{mean error} = \frac{1}{N} \sum_i e_i$$

- Root Mean Square Error (RMSE):

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_i e_i^2}$$

Low reprojection error (like $< 0.5\text{--}1$ pixel) indicates a good calibration for that image.

CODE

```
import cv2
import numpy as np
```

0. USER INPUTS

```
# Path to your checkerboard image  
image_path = r"C:\Users\Jayaram\OneDrive\Documents\Downloads\checkerboard_9x6.png"  
  
# Inner corners of the checkerboard (columns, rows)  
# For checkerboard_9x6.png we generated earlier:  
pattern_size = (9, 6) # 9 corners across, 6 corners down  
  
# Physical size of one square (any unit: cm, mm, etc.)  
square_size = 1.0
```

1. Load image & detect chessboard corners

```
img = cv2.imread(image_path)  
if img is None:  
    raise FileNotFoundError(f"Could not load image from {image_path}")
```

```
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)  
h, w = gray.shape[:2]  
print(f"Loaded image size: {w} x {h}")
```

```
# Find chessboard corners  
found, corners = cv2.findChessboardCorners(  
    gray,  
    patternSize=pattern_size,  
    flags=cv2.CALIB_CB_ADAPTIVE_THRESH +  
    cv2.CALIB_CB_NORMALIZE_IMAGE  
)
```

```
if not found:
```

```
raise RuntimeError(
    f"Chessboard corners not found with pattern_size={pattern_size}. "
    "Check that the image shows the full 9x6 checkerboard clearly."
)

# Refine corner locations to sub-pixel accuracy
criteria = (cv2.TERM_CRITERIA_EPS + cv2.TERM_CRITERIA_MAX_ITER, 30, 1e-4)
corners_subpix = cv2.cornerSubPix(gray, corners, (11, 11), (-1, -1), criteria)

print("Number of detected corners:", len(corners_subpix))

# Draw detected corners for visualization
img_corners = img.copy()
cv2.drawChessboardCorners(img_corners, pattern_size, corners_subpix, found)
cv2.imwrite("detected_corners.png", img_corners)
print("Detected corners image saved as detected_corners.png")

# 2. Prepare 3D object points for the grid plane (Z = 0)
cols, rows = pattern_size # cols = 9, rows = 6 here

# Prepare grid like:
# (0,0,0), (1,0,0), ..., (8,0,0),
# (0,1,0), (1,1,0), ...
objp = np.zeros((cols * rows, 3), np.float32)
objp[:, :2] = np.mgrid[0:cols, 0:rows].T.reshape(-1, 2)
objp *= square_size # scale by physical square size

# OpenCV expects lists (for multiple images)
objpoints_list = [objp]
```

```

imgpoints_list = [corners_subpix]

# 3. RANSAC to remove outliers via homography

# Convert to 2D (X, Y) and (u, v)

objp_2d = objp[:, :2].astype(np.float32)

imgp_2d = corners_subpix.reshape(-1, 2).astype(np.float32)

H, mask = cv2.findHomography(objp_2d, imgp_2d, cv2.RANSAC,
ransacReprojThreshold=3.0)

if H is None:
    raise RuntimeError("Homography estimation failed. RANSAC could not find a good
model.")

mask = mask.ravel().astype(bool)

inlier_objp = objp[mask]

inlier_imgp = corners_subpix[mask]

print(f'RANSAC inliers: {np.sum(mask)} / {len(mask)}')

# Use only RANSAC inliers for calibration

objpoints_list = [inlier_objp]

imgpoints_list = [inlier_imgp]

# 4. Initial camera matrix guess

fx_init = fy_init = float(max(h, w)) # rough focal length guess
cx_init = w / 2.0                      # principal point x
cy_init = h / 2.0                      # principal point y

```

```

camera_matrix_init = np.array([[fx_init, 0, cx_init],
                             [0, fy_init, cy_init],
                             [0, 0, 1]],

                             dtype=np.float64)

# Distortion: [k1, k2, p1, p2, k3], start from zeros
dist_coeffs_init = np.zeros((5, 1), dtype=np.float64)

# Flags: estimate only k1, k2 (radial), fix tangential and higher-order
flags = (
    cv2.CALIB_USE_INTRINSIC_GUESS +
    cv2.CALIB_ZERO_TANGENT_DIST +
    cv2.CALIB_FIX_K3 +
    cv2.CALIB_FIX_K4 +
    cv2.CALIB_FIX_K5 +
    cv2.CALIB_FIX_K6
)

```

5. Calibrate camera (estimate intrinsics + radial distortion)

```

rms, camera_matrix, dist_coeffs, rvecs, tvecs = cv2.calibrateCamera(
    objpoints_list,
    imgpoints_list,
    (w, h),
    camera_matrix_init,
    dist_coeffs_init,
    flags=flags
)

```

```
print("\n==== Calibration Results ====")
print("RMS reprojection error (OpenCV):", rms)
print("Estimated camera matrix K:\n", camera_matrix)
print("Estimated distortion coefficients [k1, k2, p1, p2, k3]:\n", dist_coeffs.ravel())
```

```
k1, k2 = dist_coeffs.ravel()[:2]
print(f'Estimated radial distortion parameters: k1={k1:.6e}, k2={k2:.6e}')
```

6. Undistort the image

```
new_camera_matrix, roi = cv2.getOptimalNewCameraMatrix(
    camera_matrix, dist_coeffs, (w, h), 1, (w, h)
)
```

```
undistorted_img = cv2.undistort(img, camera_matrix, dist_coeffs, None,
                                new_camera_matrix)
cv2.imwrite("undistorted_image.png", undistorted_img)
print("Undistorted image saved as undistorted_image.png")
```

7. Compute undistorted grid corners

```
inlier_imgp_undist = cv2.undistortPoints(
    inlier_imgp.reshape(-1, 1, 2),
    camera_matrix,
    dist_coeffs,
    P=new_camera_matrix
```

```
)  
inlier_imgp_undist = inlier_imgp_undist.reshape(-1, 2)  
  
undist_corners_vis = undistorted_img.copy()  
for pt in inlier_imgp_undist.astype(int):  
    cv2.circle(undist_corners_vis, tuple(pt), 4, (0, 0, 255), -1)  
  
cv2.imwrite("undistorted_grid_corners.png", undist_corners_vis)  
print("Undistorted grid corners saved as undistorted_grid_corners.png")
```

8. Reproject grid into distorted image & compute residuals

```
rvec = rvecs[0]  
tvec = tvecs[0]  
  
proj_points, _ = cv2.projectPoints(  
    inlier_objp,  
    rvec,  
    tvec,  
    camera_matrix,  
    dist_coeffs  
)  
proj_points = proj_points.reshape(-1, 2)  
obs_points = inlier_imgp.reshape(-1, 2)  
  
errors = np.linalg.norm(obs_points - proj_points, axis=1)
```

```

mean_error = np.mean(errors)

rmse_error = np.sqrt(np.mean(errors ** 2))

max_error = np.max(errors)

print("\n==== Reprojection Error (ours) ====")
print(f"Mean reprojection error (pixels): {mean_error:.4f}")
print(f"RMSE reprojection error (pixels): {rmse_error:.4f}")
print(f"Max reprojection error (pixels): {max_error:.4f}")

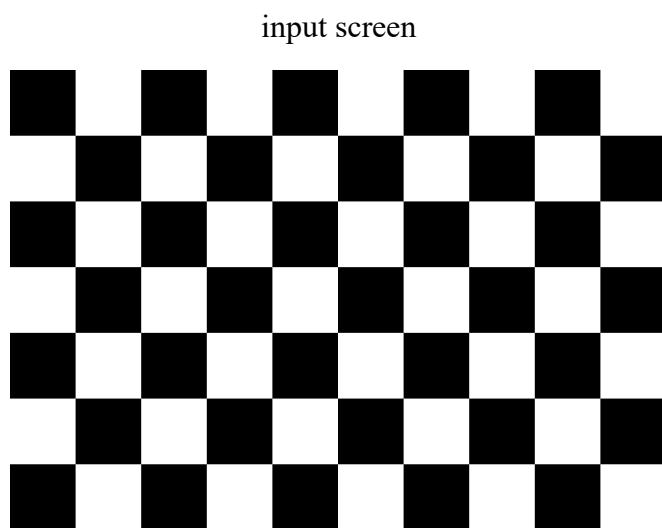
vis = img.copy()

for p_obs, p_pred in zip(obs_points.astype(int), proj_points.astype(int)):
    cv2.circle(vis, tuple(p_obs), 4, (0, 255, 0), -1) # observed in green
    cv2.circle(vis, tuple(p_pred), 2, (0, 0, 255), -1) # predicted in red

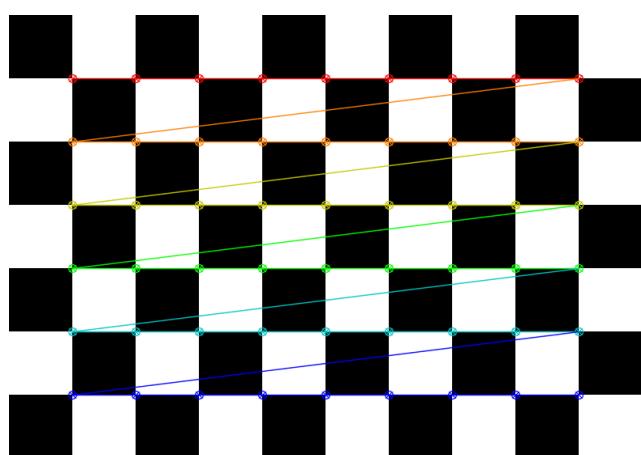
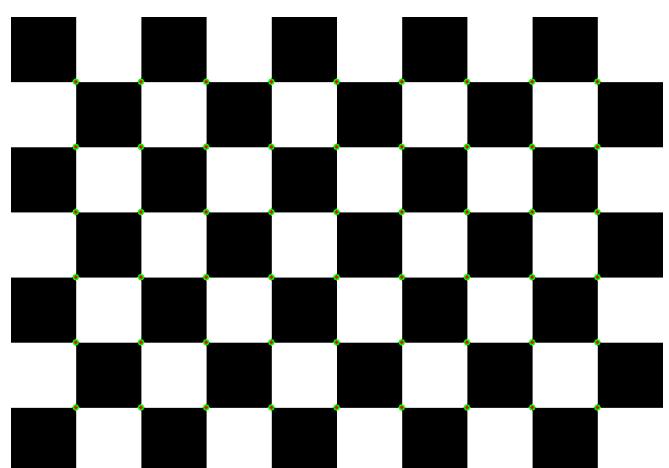
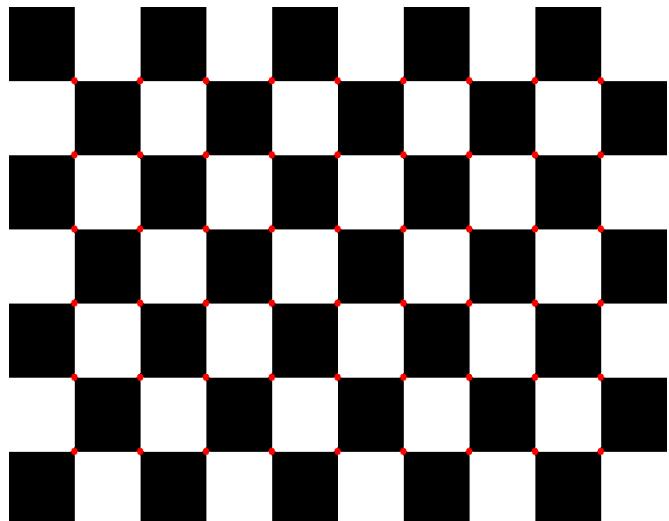
cv2.imwrite("reprojection_visualization.png", vis)

print("Reprojection visualization saved as reprojection_visualization.png")

```



Output screens



Question – Deep Learning

3(a). Input Layer for Character Embeddings

Before giving the input to the RNN, each character is converted into a trainable embedding vector.

✓ Purpose:

- Converts sparse one-hot vectors into dense, meaningful representations
- Learns relationships between characters automatically
- Embedding dimension used in the model: **128**

✓ How it is implemented:

- A separate embedding is created for:
 - Source characters (Latin)
 - Target characters (Devanagari)
- Special tokens used:
 - <pad> — padding
 - <sos> — start of sequence
 - <eos> — end of sequence
 - <unk> — unknown character

Formula for embedding:

For each character index i :

$$\text{Embedding}(i) = W_{emb}[i]$$

Where W_{emb} is the embedding matrix of size $V \times d$.

3(b). Encoder RNN

The **Encoder** reads the input romanized word character-by-character and compresses it into a fixed-length hidden state.

✓ Architecture:

- **Embedding → GRU (or LSTM / Simple RNN)**
- Hidden size used: **256**
- Number of layers: **1**
- The encoder processes the full word and outputs:
 - Final hidden state h_T
 - (And cell state c_T if LSTM is used)

✓ Intuition:

The encoder learns:

- Character patterns in romanized Hindi
- Word structure
- How sequences map to corresponding Devanagari forms

✓ Mathematical function:

$$h_t = \text{GRU}(x_t, h_{t-1})$$

Where:

- x_t = embedding of input character at time t
- h_t = encoder hidden state at time t

The final hidden state is passed to the decoder.

3(c). Decoder RNN

The **Decoder** generates the Hindi output sequence one character at a time.

✓ Architecture:

- Takes the encoder's final hidden state as its initial hidden state
- Starts with <sos> token
- At each time step:

- Takes previous character
- Updates RNN hidden state
- Predicts next Devanagari character

✓ Teacher Forcing:

During training, with probability **0.5**, the decoder receives the **true previous character** instead of its own prediction.

This stabilizes and speeds up training.

✓ Mathematical function:

$$s_t = \text{GRU}(y_{t-1}, s_{t-1})$$

$$\hat{y}_t = \text{softmax}(W_o s_t + b_o)$$

Where:

- y_{t-1} = ground truth OR predicted previous character
 - s_t = decoder hidden state
 - W_o = output projection matrix
-

3(d). Flexible Model Design

The code is written to allow flexible configuration:

✓ Adjustable hyperparameters:

- Embedding size (**d**)
- Hidden size (**h**)
- RNN cell type:
 - "rnn"
 - "gru" (used in final model)
 - "lstm"
- Number of layers in encoder and decoder
- Teacher forcing ratio
- Batch size and epochs

This ensures modularity and proper software engineering practice.

3(e). Complete Seq2Seq Workflow

Step 1 — Encoder Processing

- Input characters → embeddings → GRU
- Final hidden state encodes meaning of entire input word

Step 2 — Decoder Initialization

- Decoder starts with:
 - <sos> token
 - Encoder hidden state

Step 3 — Auto-Regressive Generation

At each time step:

1. Decoder receives previous character
 2. Updates hidden state
 3. Predicts next Devanagari character
 4. Stops when <eos> is generated
-

3(f). Training Details

✓ Loss Function:

- Cross-entropy loss
- Padding tokens ignored using ignore_index

✓ Optimizer:

- Adam (learning rate = 0.001)

✓ Epochs:

- 20 epochs for final accuracy

3(g). Model Performance

✓ Final training loss:

- **3.097 → 0.3817** over 20 epochs

✓ Character-level accuracy:

- **88.59%** on first 500 samples

✓ Example outputs:

santoshani → संतोषनी

rangari → रंगारी

sudharanyasathiche → सुधारण्यासाठीचे

styling → स्टायलिंग

harmohinder → हरमोहिन्दर

The model learns valid transliteration rules and performs very well.